

Project 3

Report

2017314331

김도엽

신동균 교수님

운영체제

0. mmap(), munmap(), freemem() 3가지 syscall 과 pagefault_handler의 구현

- 1) 모든 syscall에 사용된 함수는 pdf에서 설명하는 모든 case를 handle 하도록 구현되었습니다.
- 2) 코드의 설명은 코드 각 줄에 존재하는 주석을 통해 진행되었습니다.

1. 3가지 syscall을 추가하기 위한 기본 setting. 알파벳 순으로 각 파일에 추가 및 수정된 코드를 설명하겠습니다.

- 1) def.h에 관련 함수 prototype 추가.

```
// file.c
Int read_from_file_area(struct file *f, int offset, char *mem, int size);

// proc.c
int      add_mmap_area(struct proc *p, uint addr, int length, int prot,
int flags);
//sys_mmap.c
uint      mmap(uint, int, int, int, int, int);
int      munmap(uint);
int      freemem(void);

// vm.c
pte_t *   walkpgdir(pte_t *pgdir, const void *va, int alloc);
int      mappages(pte_t *pgdir, void *va, uint size, uint pa, int perm);
```

- 2) file.h에 file을 읽어오기 위한 함수 추가.

```
int read_from_file_area(struct file *f, int offset, char *mem, int size) {
    int original_offset = f->off;
    f->off = offset;
    int result = fileread(f, mem, size);
    f->off = original_offset;
    return result;
}
```

- 3) 기존의 file.h의 헤더파일을 전방선언하지 못해 컴파일 문제가 발생. 이 문제를 해결하기 위한 file1.h 을 만들어 file.h에 존재하는 struct file 구조체를 선언.

```
// file1.h
struct file {
    enum { FD_NONE, FD_PIPE, FD_INODE } type;
    int ref;
    char readable;
    char writable;
    struct pipe *pipe;
```

```
struct inode *ip;
uint off;
```

4) kalloc.c에 freemem 함수에서 쓰일 freed page를 세는 전역변수 선언.
memory allocation과 free가 일어날 때 해당 변수를 센다.

```
void
kfree(char *v)
{
    //..기존 field
    free_page_count++; // 추가
}
void
kalloc(void)
{
    //..기존 field
    free_page_count--; // 추가
}
```

5) param.h에 prot 변수, mmaping memory flag 변수 및 mmapbase 선언.

```
#define PROT_READ 0x1
#define PROT_WRITE 0x2
#define MAP_ANONYMOUS 0x1
#define MAP_POPULATE 0x2
#define MMAPBASE 0x40000000
#define MAX_MMAP_AREA 64
```

6) proc.c에 mmap에 필요한 함수 선언.

6-1) find_mmap_area : process와 page address를 이용하여 mmap_area를 search하는 함수.

```
struct mmap_area *find_mmap_area(struct proc *p, uint addr) {
    uint pg_addr = PGROUNDDOWN(addr);
    for (int i = 0; i < MAX_MMAP_AREA; i++) {
        struct mmap_area *area = &p->mmap_areas[i];
        if (area->addr <= pg_addr && pg_addr < area->addr + area->length) {
            return area;
        }
    }
    return 0;
}
```

6-2) allocproc() 에서 mmap_area initializing 진행.

```
static struct proc*
allocproc(void)
{
    //.. written codes
    for (int i = 0; i < MAX_MMAP_AREA; i++) {
        p->mmap_areas[i].f = 0;
        p->mmap_areas[i].addr = 0;
        p->mmap_areas[i].length = 0;
        p->mmap_areas[i].offset = 0;
        p->mmap_areas[i].prot = 0;
        p->mmap_areas[i].flags = 0;
        p->mmap_areas[i].p = 0;
    }
    //..written codes
}
```

6-3) fork() 에서 자식 프로세스가 부모 프로세스의 mmap 을 복제

```
int
fork(void)
{
    memmove(np->mmap_areas, curproc->mmap_areas, sizeof(curproc->mmap_areas));
    for (int i = 0; i < MAX_MMAP_AREA; i++) {
        if (np->mmap_areas[i].addr != 0) {
            np->mmap_areas[i].p = np;
            if (np->mmap_areas[i].f)
                filedup(np->mmap_areas[i].f);
        }
    }
}
```

7) proc.h에서 mmap area array 선언.

```
struct mmap_area {
    struct file *f;
    uint addr;
    int length;
    int offset;
    int prot;
    int flags;
    struct proc *p; // the process with this mmap_area
};
struct proc {
    //..written codes
    struct mmap_area mmap_areas[MAX_MMAP_AREA];
}
```

8) syscall.c 에서의 syscall list 삽입 및 syscall 함수 prototype 선언.

```
extern int sys_mmap(void);
extern int sys_munmap(void);
extern int sys_freemem(void);

static int (*syscalls[])(void) = {
[SYS_mmap]      = sys_mmap,
[SYS_munmap]    = sys_munmap,
[SYS_freemem]   = sys_freemem,
}
```

9) syscall.h syscall number 추가.

```
#define SYS_mmap 26
#define SYS_munmap 27
#define SYS_freemem 28
```

10) sysporc.c 에서 wrapper syscall 함수 등록을 위한 sys_mmap, sys_munmap, sys_freemem 함수 선언.

```
int
sys_mmap(void)
{
    //uint addr;
    int addr, length, prot, flags, fd, offset;

    if (argint(0, &addr) < 0 ||
        argint(1, &length) < 0 ||
        argint(2, &prot) < 0 ||
        argint(3, &flags) < 0 ||
        argint(4, &fd) < 0 ||
        argint(5, &offset) < 0)
        return -1;

    return (int)mmap(addr, length, prot, flags, fd, offset);
}

int
sys_munmap(void)
{
    uint addr;
    if (argint(0, (int *)&addr) < 0)
        return -1;
}
```

```

    return munmap(addr);
}

int
sys_freemem(void)
{
    return freemem();
}

```

11) user.h 에 sys_call에 쓰일 실제 구현된 mmap(), munmap() , freemem() 함수 prototype 전역 선언.

```

uint mmap(uint addr, int length, int prot, int flags, int fd, int offset);
int munmap(uint addr);
int freemem(void);

```

12) usys.S에 SYSCALL 매크로 함수 선언

```

SYSCALL(mmap)
SYSCALL(munmap)
SYSCALL(freemem)

```

2. sys_mmap.c라는 이름의 별도의 파일에 mmap, munmap, freemem 함수 선언 및 작성.

1) mmap() 함수 선언

```

uint
mmap(uint addr, int length, int prot, int flags, int fd, int offset)
{
    struct proc *p = myproc();
    struct mmap_area *area = 0;
    struct file *f = 0;

    // Check that addr is page aligned and length is positive
    if (addr % PGSIZE != 0 || length <= 0)
        return -1;
}

```

```

if ((flags & MAP_ANONYMOUS) == 0) {
    if (fd == -1)
        return -1; // Invalid file descriptor

    f = p->ofile[fd];
    if (!f)
        return -1; // Invalid file descriptor

    if ((prot & PROT_WRITE) && (f->writable == 0))
        return -1; // Protection mismatch with file's open flags

    filedup(f); // increase the reference count of the file
}

// Check for overlapping mappings
for (int i = 0; i < MAX_MMAP_AREA; i++) {
    if (p->mmap_areas[i].addr != 0 &&
        ((p->mmap_areas[i].addr <= addr && p->mmap_areas[i].addr + p-
>mmap_areas[i].length > addr) ||
        (p->mmap_areas[i].addr >= addr && addr + length > p-
>mmap_areas[i].addr))) {
        return -1; // Overlapping mapping detected
    }
}

// Find an available mmap_area
for (int i = 0; i < MAX_MMAP_AREA; i++) {
    if (p->mmap_areas[i].addr == 0) {
        area = &p->mmap_areas[i];
        break;
    }
}

if (!area)
    return -1; // No free mmap_area available
area->f = f;
area->addr = MMAPBASE + addr;
area->length = length;
area->offset = offset;
area->prot = prot;
area->flags = flags;
area->p = p;

if (flags & MAP_POPULATE) {
    // Allocate physical pages and create the page table entries for the
entire mapping area

```

```

    for (uint a = area->addr; a < area->addr + area->length; a += PGSIZE)
{
    char *mem = kalloc();
    cprintf("mmap called with address: %x, length: %d\n", addr,
length);
    if (!mem) {
        // Handle error: failed to allocate memory
        return -1;
    }

    if(walkpgdir(p->pgdir, (void*)a, 0) != 0) {
        cprintf(">> Address already mapped: %x\n", a);
        continue; // Skip this address
    }
    if(mappages(p->pgdir, (void*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) <
0) {
        // Handle error: failed to map pages
        kfree(mem); // Don't forget to free the allocated page in
case of error
        return -1;
    }
}
else if (flags & MAP_ANONYMOUS) {
    // For anonymous mapping, we still need to allocate physical pages and
create the page table entries
    for (uint a = area->addr; a < area->addr + area->length; a += PGSIZE)
{
        char *mem = kalloc();
        if (!mem) {
            // Handle error: failed to allocate memory
            return -1;
        }
        memset(mem, 0, PGSIZE);
        if (mappages(p->pgdir, (void*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) <
0) {
            // Handle error: failed to map pages
            cprintf(">> Handle error: failed to map pages\n");
            kfree(mem);
            // Don't forget to free the allocated page in case of error
            return -1;
        }
    }
}
return area->addr;
}

```


2) munmap() 함수 구현을 위한 내부 함수 및 munmap() 함수 :

```
void
freepgtable(pde_t *pgdir, void *va, uint size)
{
    char *a, *last;
    pte_t *pte;

    a = (char*)PGROUNDDOWN((uint)va);
    last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
    for(; a <= last; a += PGSIZE){
        if((pte = walkpgdir(pgdir, a, 0)) == 0)
            continue;
        if(*pte & PTE_P){
            char *v = P2V(PTE_ADDR(*pte));
            kfree(v);
        }
        *pte = 0;
    }
}

int
remove_mmap_area(struct proc *p, uint addr)
{
    int found = -1;

    // Traverse the array of mmap_areas
    for (int i = 0; i < MAX_MMAP_AREA; i++) {
        if (p->mmap_areas[i].addr == addr) { // If the current area is the
target area
            found = i;
            break;
        }
    }

    // If the target area wasn't found, return -1
    if (found == -1)
        return -1;

    // Mark the found mmap_area as unused. Here I'm assuming that an addr of 0
means the area is unused.
    p->mmap_areas[found].addr = 0;
    p->mmap_areas[found].length = 0;
    p->mmap_areas[found].flags = 0;
    p->mmap_areas[found].prot = 0;
    p->mmap_areas[found].offset = 0;
    p->mmap_areas[found].f = 0;
}
```

```

    p->mmap_areas[found].p = 0;

    return 1;
}

int
munmap(uint addr)
{
    struct proc *p = myproc();
    struct mmap_area *area = find_mmap_area(p, addr);
    if (!area)
        return -1;

    pte_t *pte;
    char *a;
    uint pa;

    for (a = (char*)PGROUNDDOWN((uint)area->addr); a < (char*)(area->addr +
area->length); a += PGSIZE) {
        pte = walkpgdir(p->pgdir, a, 0);
        if (!pte)
            continue;
        if (*pte & PTE_P) {
            pa = PTE_ADDR(*pte);
            if(pa == 0)
                panic("kfree");
            char *v = P2V(pa);
            memset(v, 1, PGSIZE);
            kfree(v);
            *pte = 0;
        }
    }

    // Removing corresponding mmap_area structure
    if(remove_mmap_area(p, addr) < 0)
        return -1;

    return 1;
}

```

3) freemem() 함수 : kalloc에 선언되어 있는 free_page_count 변수를 return.

```

int
freemem(void)
{
    extern int free_page_count; // Import the global variable from kalloc.c
    return free_page_count;
}

```